KQL Cheat Sheet for Real Time Intelligence

KQL

A comprehensive reference for Kusto Query Language (KQL) specifically tailored for Real Time Intelligence scenarios.



© KQL Query Flow

```
graph TD
   A[Data Source] --> B[Filter]
    B --> C[Transform]
   C --> D[Aggregate]
   D --> E[Present]
    B1[where] --> B
    B2[take] --> B
    C1[project] --> C
    C2[extend] --> C
    C3[parse] --> C
    D1[summarize] --> D
    D2[count] --> D
    D3[distinct] --> D
    E1[sort] --> E
    E2[top] --> E
    E3[render] --> E
    style A fill:#e1f5fe
    style B fill:#f3e5f5
    style C fill:#e8f5e8
    style D fill:#fff3e0
    style E fill:#fce4ec
```

Table of Contents

- Basic Syntax
- <u>Data Types</u>
- Common Operators
- <u>Time Functions</u>
- Aggregations
- <u>Joins</u>
- Advanced Patterns

- Performance Tips
- Real Time Intelligence Specific
- Security & Threat Hunting
- Community Queries

Basic Syntax

Table Selection and Basic Filtering

Documentation: where operator | Query fundamentals

```
// Basic table query
TableName

// Filter by specific condition
TableName
| where ColumnName == "value"

// Multiple conditions
TableName
| where ColumnName == "value" and TimeGenerated > ago(1h)

// Case-insensitive string comparison
TableName
| where ColumnName =~ "value"
```

Common Projections

Documentation: project operator extend operator

```
// Select specific columns
TableName
| project TimeGenerated, ColumnName, AnotherColumn

// Rename columns
TableName
| project Timestamp = TimeGenerated, Name = ColumnName

// Create calculated columns
TableName
| project TimeGenerated, Duration = EndTime - StartTime
```

Data Types

```
mindmap
root((KQL Data Types))
String
string
```

```
dynamic
Numeric
int
long
real
decimal
DateTime
datetime
timespan
Boolean
bool
Special
guid
dynamic
```

String Operations

Documentation: String operators | extract() function | split() function

```
// Contains (case-insensitive)
| where Message contains "error"

// Starts with
| where EventName startswith "Microsoft"

// Regular expressions
| where Message matches regex @"Error \d+"

// String extraction
| extend ErrorCode = extract(@"Error (\d+)", 1, Message)

// String splitting
| extend Parts = split(Message, ";")
```

Numeric Operations

```
// Basic math
| extend Total = Quantity * Price

// Rounding
| extend RoundedValue = round(Value, 2)

// Absolute value
| extend AbsValue = abs(Value)
```

DateTime Operations

```
// Current time
| extend Now = now()

// Time ago
| where TimeGenerated > ago(1h)
```

```
| where TimeGenerated > ago(7d)

// Date formatting
| extend FormattedDate = format_datetime(TimeGenerated, "yyyy-MM-dd HH:mm")

// Date parts
| extend Hour = datetime_part("hour", TimeGenerated)
| extend DayOfWeek = dayofweek(TimeGenerated)
```

Common Operators

Filtering

```
// In operator
| where EventLevel in ("Error", "Warning")

// Between
| where ResponseTime between (100 .. 500)

// Not null
| where isnotnull(ColumnName)

// Empty or null
| where isempty(ColumnName) or isnull(ColumnName)
```

Sorting and Limiting

```
// Sort ascending
| sort by TimeGenerated asc

// Sort descending (default)
| sort by Count desc

// Top N records
| top 10 by Count desc

// Skip and take (pagination)
| sort by TimeGenerated desc
| serialize rn = row_number()
| where rn > 10 and rn <= 20</pre>
```

Time Functions

Time Ranges

Documentation: ago() function | bin() function | Datetime/timespan arithmetic

```
// Last hour
| where TimeGenerated > ago(1h)

// Between specific times
| where TimeGenerated between (datetime(2024-01-01) .. datetime(2024-01-02))

// Time bins
| summarize Count = count() by bin(TimeGenerated, 5m)

// Start of day/week/month
| extend StartOfDay = startofday(TimeGenerated)
| extend StartOfWeek = startofweek(TimeGenerated)
| extend StartOfMonth = startofmonth(TimeGenerated)
```

Aggregations

Basic Aggregations

```
// Count all records
| summarize count()
// Count with grouping
| summarize Count = count() by EventLevel
// Multiple aggregations
summarize
   TotalEvents = count(),
   UniqueUsers = dcount(UserId),
   AvgDuration = avg(Duration),
   MaxDuration = max(Duration)
by EventLevel
// Percentiles
summarize
   P50 = percentile(Duration, 50),
   P95 = percentile(Duration, 95),
   P99 = percentile(Duration, 99)
```

Advanced Aggregations

```
// Array aggregation
| summarize EventTypes = make_set(EventType) by UserId

// List with duplicates
| summarize EventList = make_list(EventType) by UserId

// String concatenation
| summarize ErrorMessages = strcat_array(make_list(Message), "; ")

// Conditional aggregation
| summarize
```

```
ErrorCount = countif(EventLevel == "Error"),
WarningCount = countif(EventLevel == "Warning")
```

Joins

Documentation: join operator | Join flavors | Join best practices

```
graph LR
   A[Table A] --> C{Join Type}
   B[Table B] --> C

C -->|inner| D[Inner Join<br/>br/>Only matching records]
C -->|leftouter| E[Left Join<br/>br/>All from A + matches from B]
C -->|rightouter| F[Right Join<br/>All from B + matches from A]
C -->|fullouter| G[Full Join<br/>All records from both]

style A fill:#e3f2fd
style B fill:#e8f5e8
style D fill:#ff3e0
style E fill:#f3e5f5
style F fill:#e0f2f1
style G fill:#fce4ec
```

Inner Join

```
Table1
| join kind=inner (
    Table2
    | project UserId, UserName
) on UserId
```

Left Join

Join with Time Window

```
Events
| join kind=inner (
    Metrics
```

```
| where TimeGenerated > ago(1h)
) on $left.TimeGenerated == $right.TimeGenerated
```

Advanced Patterns

Window Functions

```
// Running total
| sort by TimeGenerated asc
| serialize RunningTotal = row_cumsum(Count)

// Previous value
| sort by TimeGenerated asc
| serialize PrevValue = prev(Count, 1)

// Rank
| sort by Count desc
| serialize Rank = row_number()
```

Dynamic Objects

```
// Parse JSON
| extend ParsedData = parse_json(JsonColumn)
| extend Value = ParsedData.property

// Create dynamic object
| extend Details = pack("Count", Count, "Timestamp", TimeGenerated)

// Array operations
| mv-expand ArrayColumn
| where ArrayColumn.property == "value"
```

Case Statements

```
| extend Severity = case(
    EventLevel == "Error", "High",
    EventLevel == "Warning", "Medium",
    EventLevel == "Information", "Low",
    "Unknown"
)
```

Graph Operators & Network Analysis

Documentation: graph operators graph-match operator Graph scenarios

Graph Data Model

```
graph LR
   A[Source Node] --> | Edge/Relationship | B[Target Node]
    A --> C[Target Node 2]
   B --> D[Target Node 3]
   C --> D
    style A fill:#e3f2fd
    style B fill:#e8f5e8
    style C fill:#fff3e0
    style D fill:#f3e5f5
```

Basic Graph Operations

Output Creating Graph from Tabular Data

```
// Convert network logs to graph structure
NetworkLogs
| where TimeGenerated > ago(1h)
project SourceIP, DestinationIP, Port, Protocol, BytesSent
extend EdgeType = strcat(Protocol, ":", Port)
// SourceIP and DestinationIP become nodes
// EdgeType defines the relationship
```

© Graph-Match Pattern Matching

```
// Find communication patterns
let GraphData = NetworkLogs
| where TimeGenerated > ago(24h)
| project SourceIP, DestinationIP, Protocol;
| graph-match (source)-[connection]->(destination)
 where source. SourceIP startswith "10.0"
  and destination.DestinationIP !startswith "10.0"
| project SourceInternal=source.SourceIP,
          {\tt DestinationExternal=destination.DestinationIP,}
          Protocol=connection.Protocol
 summarize Connections=count() by SourceInternal, DestinationExternal
 top 10 by Connections
```

Advanced Graph Analysis Patterns

Multi-Hop Path Analysis

```
// Find 2-hop communication paths (A -> B -> C)
let NetworkGraph = NetworkLogs
| where TimeGenerated > ago(6h)
project Source=SourceIP, Target=DestinationIP, Protocol, Timestamp=TimeGenerated;
NetworkGraph
graph-match (node1)-[edge1]->(node2)-[edge2]->(node3)
 where edge1.Timestamp < edge2.Timestamp \ \ // Ensure temporal order
 and datetime_diff('minute', edge2.Timestamp, edge1.Timestamp) <= 30 // Within 30 minutes
l project
   Path = strcat(node1.Source, " -> ", node2.Source, " -> ", node3.Target),
   FirstHop = edge1.Protocol,
   SecondHop = edge2.Protocol,
   Duration = datetime_diff('minute', edge2.Timestamp, edge1.Timestamp)
| summarize PathCount = count() by Path, FirstHop, SecondHop
sort by PathCount desc
```

Circular Communication Detection

```
// Detect circular communication patterns (A -> B -> A)
let CommunicationGraph = NetworkLogs
 where TimeGenerated > ago(2h)
| project SourceIP, DestinationIP, TimeGenerated;
CommunicationGraph
graph-match (nodeA)-[forward]->(nodeB)-[backward]->(nodeC)
 where nodeA.SourceIP == nodeC.DestinationIP // Same starting and ending node
 and nodeB.DestinationIP == nodeC.SourceIP
                                              // Reverse path
 and forward.TimeGenerated < backward.TimeGenerated
project
   CircularPath = strcat(nodeA.SourceIP, " <-> ", nodeB.DestinationIP),
   ForwardTime = forward.TimeGenerated,
   BackwardTime = backward.TimeGenerated,
   ResponseTime = datetime_diff('second', backward.TimeGenerated, forward.TimeGenerated)
| where ResponseTime <= 300 // Within 5 minutes
summarize
   CircularConnections = count(),
   AvgResponseTime = avg(ResponseTime)
by CircularPath
| sort by CircularConnections desc
```

Security Analysis with Graph Operators

Lateral Movement Detection

```
// Detect potential lateral movement in network
let AuthEvents = SecurityEvent
| where TimeGenerated > ago(24h)
 where EventID in (4624, 4625) // Logon events
| project TimeGenerated, Account, Computer, LogonType, EventID;
\ensuremath{//} Find accounts that accessed multiple computers
| graph-match (user)-[logon1]->(computer1), (user)-[logon2]->(computer2)
  where computer1.Computer != computer2.Computer // Different computers
  and logon1.EventID == 4624 and logon2.EventID == 4624 // Successful logons
  and datetime_diff('hour', logon2.TimeGenerated, logon1.TimeGenerated) between (0 .. 6)
project
```

```
SuspiciousAccount = user.Account,
   Computer1 = computer1.Computer,
   Computer2 = computer2.Computer,
   TimeSpan = datetime_diff('minute', logon2.TimeGenerated, logon1.TimeGenerated)
| summarize
   ComputersAccessed = dcount(Computer2),
   AccessPattern = make_list(Computer2)
by SuspiciousAccount
| where ComputersAccessed >= 3 // Accessed 3+ different computers
| project SuspiciousAccount, ComputersAccessed, AccessPattern
```

Suspicious Communication Chains

```
// Find suspicious communication chains (potential data exfiltration)
let SuspiciousGraph = NetworkLogs
| where TimeGenerated > ago(4h)
 where DestinationPort in (80, 443, 22, 3389) // Common ports
 project SourceIP, DestinationIP, DestinationPort, BytesSent, TimeGenerated;
SuspiciousGraph
graph-match (internal)-[connection1]->(external)-[connection2]->(final)
  where internal.SourceIP startswith "192.168" // Internal network and external.DestinationIP !startswith "192.168" // External destination
  and final.DestinationIP !startswith "192.168" // Another external destination
  and connection1.BytesSent > 100000
                                                      // Large data transfer
| project
    InternalSource = internal.SourceIP,
    FirstExternal = external.DestinationIP,
    FinalDestination = final.DestinationIP,
    DataTransferred = connection1.BytesSent,
    ConnectionTime = connection1.TimeGenerated
l summarize
    TotalDataExfiltrated = sum(DataTransferred),
    UniqueExternalTargets = dcount(FinalDestination),
    ExternalTargets = make_set(FinalDestination)
by InternalSource
| where TotalDataExfiltrated > 1000000 // More than 1MB total
| sort by TotalDataExfiltrated desc
```

III Business Intelligence with Graph Analysis

Customer Journey Mapping

```
// Map customer journey through application features
let UserEvents = customEvents
| where timestamp > ago(7d)
| where name in ("PageView", "FeatureUse", "Purchase", "Support")
| project user_Id, EventName=name, timestamp, FeatureName=tostring(customDimensions["feature"]);
UserEvents
| graph-match (user)-[event1]->(feature1)-[event2]->(feature2)
    where event1.timestamp < event2.timestamp
    and datetime_diff('hour', event2.timestamp, event1.timestamp) <= 24
| project
    UserJourney = strcat(feature1.FeatureName, " -> ", feature2.FeatureName),
        user_Id = user.user_Id,
        StepDuration = datetime_diff('minute', event2.timestamp, event1.timestamp)
| summarize
        JourneyCount = dcount(user_Id),
```

```
AvgStepDuration = avg(StepDuration)
by UserJourney
| where JourneyCount >= 10
| sort by JourneyCount desc
| project UserJourney, Users=JourneyCount, AvgDurationMins=round(AvgStepDuration, 1)
```

Influencer Network Analysis

```
// Analyze user referral patterns
let UserReferrals = customEvents
| where timestamp > ago(30d)
 where name == "UserReferral"
project
   ReferrerUser = tostring(customDimensions["referrer"]),
   NewUser = tostring(customDimensions["newUser"]),
   timestamp:
UserReferrals
graph-match (influencer)-[referral1]->(user1)-[referral2]->(user2)
 where referral1.timestamp < referral2.timestamp
   TopInfluencer = influencer.ReferrerUser,
   DirectReferral = user1.NewUser,
   IndirectReferral = user2.NewUser,
   InfluenceChain = strcat(influencer.ReferrerUser, " -> ", user1.NewUser, " -> ", user2.NewUser)
| summarize
   DirectReferrals = dcount(DirectReferral),
   IndirectReferrals = dcount(IndirectReferral),
   TotalInfluence = dcount(DirectReferral) + dcount(IndirectReferral)
by TopInfluencer
 sort by TotalInfluence desc
 take 10
| project Influencer=TopInfluencer, DirectReferrals, IndirectReferrals, TotalInfluence
```

🦴 Graph Operator Performance Tips

Optimization Strategies

Memory-Efficient Graph Queries

```
// Use summarize to reduce data size before graph operations
let SummaryGraph = NetworkLogs
 where TimeGenerated > ago(6h)
| summarize ConnectionCount = count(), TotalBytes = sum(BytesSent)
 by SourceIP, DestinationIP, Protocol // Aggregate similar connections
| where ConnectionCount >= 5; // Focus on frequent connections
SummaryGraph
graph-match (source)-[connection]->(destination)
 where connection. Total Bytes > 1000000 // Large data transfers only
| project source.SourceIP, destination.DestinationIP, connection.TotalBytes
```

© Real-World Graph Scenarios

IT Infrastructure Mapping

```
// Map service dependencies from telemetry
dependencies
 where timestamp > ago(1h)
 where success == true
 project CallingService=cloud_RoleName, CalledService=target, ResponseTime=duration
 graph-match (service1)-[dependency]->(service2)
summarize
    CallCount = count(),
    AvgResponseTime = avg(dependency.ResponseTime),
    MaxResponseTime = max(dependency.ResponseTime)
by service1.CallingService, service2.CalledService
| extend HealthStatus = case(
    AvgResponseTime < 100, "Healthy",
AvgResponseTime < 500, "Warning",</pre>
    "Critical"
| project CallingService, CalledService, CallCount,
          AvgResponseMs=round(AvgResponseTime, 1), HealthStatus
```

Fraud Detection Network

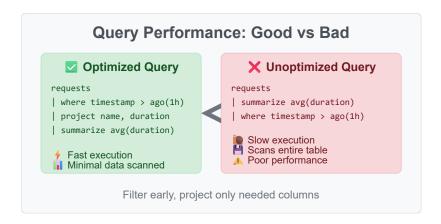
```
// Detect potentially fraudulent transaction networks
let TransactionGraph = PaymentEvents
 where timestamp > ago(24h)
| project FromAccount, ToAccount, Amount, timestamp;
TransactionGraph
| graph-match (account1)-[transfer1]->(account2)-[transfer2]->(account3)
 where transfer1.Amount == transfer2.Amount // Same amount transferred
 and datetime_diff('minute', transfer2.timestamp, transfer1.timestamp) <= 10 // Quick succession
 and transfer1. Amount > 1000 // High value transactions
project
   SuspiciousChain = strcat(account1.FromAccount, " -> ", account2.ToAccount, " -> ", account3.To
   Amount = transfer1.Amount,
   TimeGap = datetime_diff('second', transfer2.timestamp, transfer1.timestamp)
| summarize ChainCount = count() by SuspiciousChain, Amount
 where ChainCount >= 2 // Pattern repeated
| sort by Amount desc
```



***** Graph Operator Cheat Sheet

Pattern	Use Case	KQL Syntax
Simple Path	A connects to B	<pre>graph-match (a)-[edge]->(b)</pre>
Two-Hop Path	A -> B -> C	graph-match (a)-[e1]->(b)-[e2]->(c)
Circular Path	A -> B -> A	<pre>graph-match (a)-[e1]->(b)-[e2]->(c) where a.id == c.id</pre>
Common Neighbor	A -> C <- B	graph-match (a)-[e1]->(c)<-[e2]-(b)
Fan-out	One source, many targets	<pre>graph-match (source)-[edge]->(target) + summarize by source</pre>
Fan-in	Many sources, one target	graph-match (source)-[edge]->(target) + summarize by target

Performance Tips



```
graph TD
   A[Query Optimization] --> B[Filter Early]
    A --> C[Project Only Needed Columns]
    A --> D[Use Time Ranges]
   A --> E[Aggregate Before Join]
    B --> B1[where timestamp > ago(1h)]
    C --> C1[project col1, col2, col3]
    D --> D1[datetime range filters]
    E --> E1[summarize before join]
    F[Performance Killers] --> G[Full Table Scans]
    F --> H[Large Joins]
    F --> I[Complex Regex]
    F --> J[No Time Filters]
    style A fill:#e8f5e8
    style F fill:#ffebee
    style B1 fill:#e3f2fd
    style C1 fill:#e3f2fd
```

```
style D1 fill:#e3f2fd
style E1 fill:#e3f2fd
```

Filtering Early

```
// Good: Filter first
TableName
| where TimeGenerated > ago(1h)
| where EventLevel == "Error"
| summarize count() by Computer

// Bad: Filter after aggregation
TableName
| summarize count() by Computer, EventLevel
| where EventLevel == "Error"
```

Efficient Time Filtering

```
// Use specific time ranges
| where TimeGenerated between (ago(1h) .. now())

// Use time bins for large datasets
| summarize count() by bin(TimeGenerated, 1h)
```

Column Selection

```
// Project early to reduce data
TableName
| project TimeGenerated, EventLevel, Message
| where EventLevel == "Error"
```

Real Time Intelligence Specific

```
graph TB
  subgraph "Real Time Intelligence Architecture"
    A[Applications] --> B[Telemetry]
    C[Infrastructure] --> B
    D[Users] --> B

B --> E[Real Time Analytics]
    E --> F[Alerts]
    E --> G[Dashboards]
    E --> H[Queries]

F --> I[Notifications]
```

```
G --> J[Visualizations]
H --> K[Insights]
end

style A fill:#e3f2fd
style C fill:#e8f5e8
style D fill:#fff3e0
style E fill:#ff8e5f5
style F fill:#ffebee
style G fill:#e0f2f1
style H fill:#fce4ec
```

Telemetry Analysis

```
// Request analysis
requests
| where timestamp > ago(1h)
summarize
   RequestCount = count(),
   AvgDuration = avg(duration),
   P95Duration = percentile(duration, 95),
   FailureRate = round(100.0 * countif(success == false) / count(), 2)
by bin(timestamp, 5m)
render timechart
// Performance counters
performanceCounters
 where timestamp > ago(6h)
 where category == "Processor" and counter == "% Processor Time"
 summarize avg(value) by bin(timestamp, 1m), computer
| render timechart
```

Exception Analysis

```
// Exception trending
exceptions
| where timestamp > ago(24h)
| summarize Count = count() by bin(timestamp, 1h), type
| render timechart

// Top exceptions
exceptions
| where timestamp > ago(24h)
| summarize Count = count(), Sample = any(outerMessage) by type
| sort by Count desc
| take 10
```

Dependency Analysis

```
// Dependency failures
dependencies
| where timestamp > ago(1h)
```

```
| where success == false
| summarize Count = count(), Types = make_set(type) by target
| sort by Count desc
```

Custom Events

```
// Business metrics
customEvents
| where timestamp > ago(1d)
| where name == "UserAction"
| extend Action = tostring(customDimensions.action)
| summarize Count = count() by bin(timestamp, 1h), Action
| render columnchart
```

Log Analytics

```
// Error analysis
traces
| where timestamp > ago(1h)
| where severityLevel >= 3 // Error and Critical
| summarize Count = count(), Sample = any(message) by bin(timestamp, 5m)
| render timechart

// Search across logs
union traces, exceptions, requests
| where timestamp > ago(1h)
| search "specific error pattern"
| project timestamp, itemType, message
| sort by timestamp desc
```

Real-time Monitoring Queries

```
// Live metrics (use in dashboards)
requests
| where timestamp > ago(5m)
| summarize
    RequestsPerMin = count(),
    AvgResponseTime = avg(duration),
    ErrorRate = round(100.0 * countif(success == false) / count(), 2)

// Health check
heartbeat
| where TimeGenerated > ago(5m)
| summarize LastHeartbeat = max(TimeGenerated) by Computer
| extend Status = iff(LastHeartbeat < ago(2m), "Down", "Up")
| where Status == "Down"</pre>
```

Time Series Analysis Pattern

```
graph LR
   A[Time Data] --> B[bin()]
   B --> C[summarize]
   C --> D[render timechart]

A1[timestamp] --> B
   B1[bin(timestamp, 15m)] --> C
   C1[count(), avg(), etc.] --> D
   D1[Line Chart] --> E[Insights]

style A fill:#e3f2fd
style B fill:#e8f5e8
style C fill:#ff3e0
style D fill:#f3e5f5
style E fill:#fce4ec
```

Time Series Example

```
// Application performance over time
requests
| where timestamp > ago(24h)
| bin timestamp to 1h
| summarize
    RequestCount = count(),
    AvgDuration = avg(duration),
    P95Duration = percentile(duration, 95),
    ErrorRate = countif(success == false) * 100.0 / count()
| render timechart with (title="Application Performance Dashboard")
```

Alerting Queries

Common Functions Reference

Documentation: <u>KQL Function Reference</u> | <u>String Functions</u> | <u>DateTime</u> Functions

String Functions

- contains Case-insensitive substring search
- startswith / endswith String prefix/suffix check
- extract Regular expression extraction
- split Split string into array
- strlen String length
- <u>substring</u> Extract substring
- tolower / toupper Case conversion

Math Functions

- abs Absolute value
- round Round to specified decimals
- floor / ceiling Round down/up
- sqrt Square root
- pow Power function
- min / max Minimum/maximum

Date Functions

- now() Current timestamp
- ago() Time in the past
- datetime() Parse datetime
- <u>format_datetime()</u> Format timestamp
- bin() Time bucketing
- <u>datetime_part()</u> Extract date component

Aggregation Functions

- count() Count rows
- dcount() Distinct count
- <u>sum()</u> Sum values
- avg() Average
- min() / max() Minimum/maximum
- percentile() Calculate percentile
- make set() Create array of unique values
- make list() Create array with duplicates

Tips for Real Time Intelligence

- 1. Use time-based partitioning Always filter by timestamp first
- 2. Leverage materialized views For frequently accessed aggregations

- 3. Use appropriate data types Choose the right type for better performance
- 4. Monitor query performance Use query statistics to optimize
- 5. Cache frequently used results Use cached queries for dashboards
- 6. Partition your data Use appropriate partitioning strategies
- 7. Use incremental processing For large datasets, process incrementally

Quick Reference Card

Must-Know Operators

- | Pipe operator (chain operations)
- where Filter rows
- project Select/rename columns
- extend Add calculated columns
- summarize Group and aggregate
- sort Order results
- top Get top N records
- join Combine tables
- union Combine similar tables
- render Visualize results

Common Patterns

```
// Time series analysis
| summarize count() by bin(timestamp, 1h)
| render timechart

// Top N analysis
| summarize Count = count() by Category
| top 10 by Count desc

// Percentage calculation
| summarize Total = count(), Success = countif(status == "success")
| extend SuccessRate = round(100.0 * Success / Total, 2)
```

Security & Threat Hunting

Documentation: Microsoft Sentinel KQL | Advanced Security Information Model (ASIM) | Threat Hunting with KQL

```
graph TB
    subgraph "Security Analytics Pipeline"
        A[Security Logs] --> B[Normalization]
        B --> C[Detection]
        C --> D[Alerting]

E[Network Logs] --> B
```

```
F[Identity Logs] --> B
   G[Endpoint Logs] --> B

C --> H[Threat Intelligence]
   C --> I[Threat Hunting]
   C --> J[Incident Response]
end

style A fill:#ffebee
style E fill:#e8f5e8
style F fill:#e3f2fd
style G fill:#ff3e0
style G fill:#ff3e5f5
style H fill:#fce4ec
```

Common Security Tables

Authentication Analysis

```
// Failed logon attempts
SecurityEvent
 where TimeGenerated > ago(24h)
 where EventID == 4625 // Failed logon
 summarize FailedAttempts = count() by Account, Computer, IpAddress
 where FailedAttempts > 10
| sort by FailedAttempts desc
// Successful logons after failed attempts
let FailedLogons = SecurityEvent
     where TimeGenerated > ago(1h)
     where EventID == 4625
    project TimeGenerated, Account, Computer, IpAddress;
let SuccessfulLogons = SecurityEvent
     where TimeGenerated > ago(1h)
     where EventID == 4624
    | project TimeGenerated, Account, Computer, IpAddress;
FailedLogons
 join kind=inner (SuccessfulLogons) on Account, Computer
 where TimeGenerated1 < TimeGenerated
project Account, Computer, IpAddress, FailedTime=TimeGenerated, SuccessTime=TimeGenerated1
```

Network Security Monitoring

```
// Suspicious outbound connections
DeviceNetworkEvents
| where TimeGenerated > ago(1h)
| where ActionType == "ConnectionSuccess"
| where RemotePort in (22, 3389, 5985, 5986) // SSH, RDP, WinRM
| where not(RemoteIP has_any ("10.", "192.168.", "172.16.", "172.17.", "172.18.", "172.19.", "172.15.", "172.19.", "172.15.", "172.19.", "172.15.", "172.19.", "172.15.", "172.19.", "172.15.", "172.17.", "172.18.", "172.19.", "172.15.", "172.19.", "172.15.", "172.17.", "172.18.", "172.19.", "172.15.", "172.17.", "172.18.", "172.19.", "172.15.", "172.17.", "172.18.", "172.19.", "172.15.", "172.17.", "172.17.", "172.18.", "172.19.", "172.15.", "172.16.", "172.17.", "172.17.", "172.18.", "172.19.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.15.", "172.1
```

Process Execution Analysis

```
// PowerShell execution with obfuscation indicators
DeviceProcessEvents
| where TimeGenerated > ago(24h)
| where FileName =~ "powershell.exe"
| where ProcessCommandLine has_any ("-enc", "-encoded", "FromBase64String", "iex", "invoke-express | project TimeGenerated, DeviceName, AccountName, ProcessCommandLine
| sort by TimeGenerated desc

// Suspicious parent-child process relationships
DeviceProcessEvents
| where TimeGenerated > ago(24h)
| where InitiatingProcessFileName in~ ("winword.exe", "excel.exe", "powerpnt.exe", "outlook.exe")
| where FileName in~ ("powershell.exe", "cmd.exe", "wscript.exe", "cscript.exe", "regsvr32.exe")
| project TimeGenerated, DeviceName, InitiatingProcessFileName, FileName, ProcessCommandLine
```

File System Security

```
// Sensitive file access
DeviceFileEvents
| where TimeGenerated > ago(24h)
| where FileName has_any ("password", "credential", "secret", "private", "key")
| where ActionType in ("FileCreated", "FileModified", "FileRenamed")
| summarize count() by DeviceName, AccountName, FileName, FolderPath
| sort by count_ desc

// Executable files written to temp directories
DeviceFileEvents
| where TimeGenerated > ago(24h)
| where ActionType == "FileCreated"
| where FileName endswith ".exe"
| where FolderPath has_any ("temp", "tmp", "appdata\\local\\temp")
| project TimeGenerated, DeviceName, AccountName, FileName, FolderPath, SHA256
```

Threat Intelligence Integration

Community Queries

♣ Credits: These queries are contributed by the cybersecurity community. Special thanks to the following sources and contributors:

Microsoft Security Community

- Microsoft Sentinel Community Official Microsoft Sentinel repository
- KQL Search Community-driven KQL query repository by @rod-trent
- Microsoft 365 Defender Hunting Queries Official hunting queries
- Bert-JanP/Hunting-Queries-Detection-Rules Comprehensive hunting query collection

Notable Contributors

- Rod Trent (@rod-trent) KQL Search platform and extensive query contributions
- Bert-Jan Piet (@Bert-JanP) Threat hunting queries and detection rules
- Cyb3rWard0g (@Cyb3rWard0g) Advanced persistent threat hunting
- FalconForce (@FalconForceTeam) Purple team and threat hunting methodologies

Advanced Persistent Threat (APT) Hunting

```
// Living off the Land (LoL) techniques detection
// Credit: Inspired by MITRE ATT&CK framework and community research
DeviceProcessEvents
| where TimeGenerated > ago(24h)
| where FileName in~ (
    "certutil.exe", "bitsadmin.exe", "wmic.exe",
    "rundll32.exe", "regsvr32.exe", "mshta.exe",
    "installutil.exe", "regasm.exe", "regsvcs.exe"
)
| where ProcessCommandLine has_any (
    "http", "https", "ftp", "download",
    "urlcache", "verifyctl", "encode", "decode"
)
| project TimeGenerated, DeviceName, AccountName, FileName, ProcessCommandLine
| sort by TimeGenerated desc
```

Supply Chain Attack Detection

Ransomware Detection

```
// Rapid file encryption indicators
// Credit: Analysis from various ransomware incident reports
DeviceFileEvents
| where TimeGenerated > ago(1h)
| where ActionType in ("FileRenamed", "FileModified")
| where FileName has_any (".encrypt", ".locked", ".crypto", ".crypt", ".enc")
    or FileName matches regex @".*\.(jpg|pdf|doc|xls|ppt|png|gif|bmp|mp3|mp4|avi|mov|zip|rar|7z|t)
| summarize
    FileCount = count(),
    UniqueExtensions = dcount(FileName),
    FileTypes = make_set(FileName)
by DeviceName, AccountName, bin(TimeGenerated, 5m)
| where FileCount > 50 // High volume of file changes
| sort by TimeGenerated desc
```

Credential Access Hunting

```
// LSASS process access patterns
// Credit: Security research on credential dumping techniques
DeviceEvents
| where TimeGenerated > ago(24h)
| where ActionType == "ProcessAccess"
| where FileName == "lsass.exe"
| where InitiatingProcessFileName !in~ ("svchost.exe", "wininit.exe", "winlogon.exe", "csrss.exe")
| summarize count() by DeviceName, InitiatingProcessFileName, InitiatingProcessCommandLine
| sort by count_ desc
```

Command and Control (C2) Detection

```
// Beaconing behavior detection
// Credit: Network security research and threat hunting methodologies
DeviceNetworkEvents
| where TimeGenerated > ago(24h)
 where ActionType == "ConnectionSuccess"
| summarize
   ConnectionCount = count(),
   UniqueRemotePorts = dcount(RemotePort),
   AvgInterval = avg(prev(TimeGenerated) - TimeGenerated)
by DeviceName, RemoteIP, bin(TimeGenerated, 1h)
 where ConnectionCount > 10 and UniqueRemotePorts < 3
 where AvgInterval between (time(0.1s) .. time(30m)) // Regular intervals
sort by ConnectionCount desc
```

Data Exfiltration Detection

```
// Large data transfers to external destinations
// Credit: Data loss prevention research and incident analysis
DeviceNetworkEvents
 where TimeGenerated > ago(24h)
where ActionType == "NetworkConnection"
 where not(RemoteIP has_any ("10.", "192.168.", "172.")) // External IPs
 summarize TotalBytes = sum(SentBytes) by DeviceName, RemoteIP, bin(TimeGenerated, 1h)
 where TotalBytes > 100000000 // >100MB
sort by TotalBytes desc
```

Persistence Mechanism Detection

```
// Registry-based persistence techniques
// Credit: MITRE ATT&CK T1547 and community research
DeviceRegistryEvents
 where TimeGenerated > ago(24h)
 where ActionType == "RegistryValueSet"
| where RegistryKey has_any (
    "\\Run\\", "\\RunOnce\\", "\\RunServices\\",
    "\\Winlogon\\", "\\Explorer\\Run\\",
    "\\Image File Execution Options\\",
    "\\App\bar{\text{Init_DLLs}", "\\ServiceDLL"
| project TimeGenerated, DeviceName, AccountName, RegistryKey, RegistryValueName, RegistryValueDat
| sort by TimeGenerated desc
```



A Threat Hunting Resources

Official Documentation

- Microsoft Sentinel Hunting Official hunting guide
- MITRE ATT&CK Framework Threat tactics and techniques

• Microsoft 365 Defender Advanced Hunting

Community Resources

- **KQL Cafe** Interactive KQL learning platform
- Azure Sentinel Notebooks Jupyter notebooks for security analysis
- <u>Uncoder.io</u> Sigma rule to KQL converter
- <u>Sentinel ATT&CK</u> MITRE ATT&CK aligned queries

Training and Certification

- Microsoft Security Operations Analyst SC-200 certification
- **<u>KQL Pluralsight Course</u>** Comprehensive KQL training
- Microsoft Learn KQL Free learning path

⚠ **Disclaimer**: These security queries are provided for educational and defensive purposes. Always test in a controlled environment and ensure compliance with your organization's policies and applicable laws.